

CodeMentor

AI-Powered Coding Assistant

Transformer-based LLMs meet backend systems to create an intelligent coding companion. Built with Ollama, Flask, and advanced prompt engineering.



Problem Statement



Learning Curve

Beginners struggle with syntax, patterns, and debugging unfamiliar codebases.



Debugging Overhead

Hours spent identifying bugs instead of building features and learning concepts.

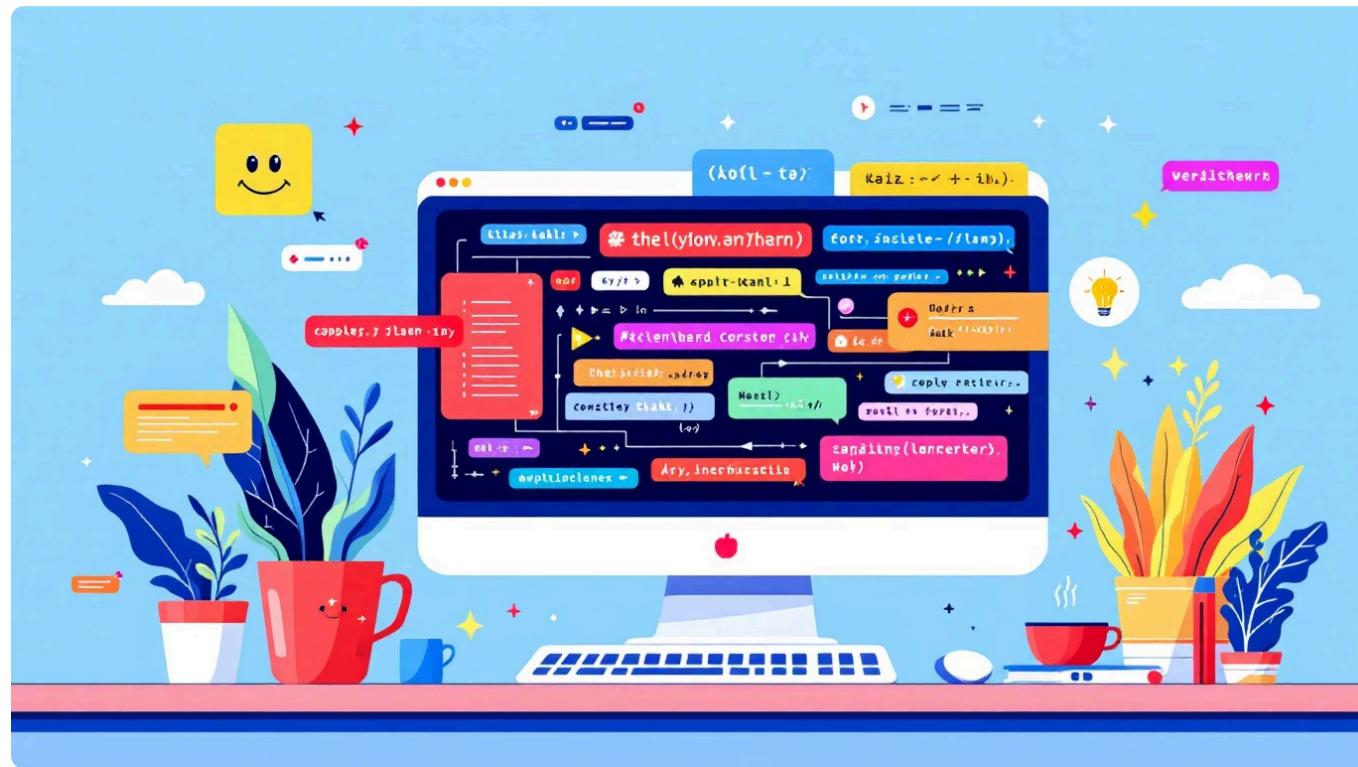


Optimization Blindspots

Difficulty identifying performance bottlenecks and refactoring opportunities in existing code.

Developers need real-time guidance that understands context, explains trade-offs, and adapts to skill levels.

Project Motivation



Democratizing Development

Lower entry barriers for students and junior developers through instant expert guidance.

Accelerated Learning

Real-time explanations transform passive coding into active learning experiences.

Productivity Multiplier

Reduce debugging time from hours to minutes with AI-powered root cause analysis.

Solution Overview



User Interface

Web-based frontend accepts code input and queries

Backend Engine

Flask API processes requests and formats prompts



LLM Core

Quantized transformer models via Ollama

Intelligent Response

Context-aware code analysis and explanations

Transformer-based architecture combines tokenization, self-attention mechanisms, and next-token prediction for contextual code understanding.

System Architecture

01

Client Request

Frontend sends code and context to Flask backend

02

Prompt Engineering

Backend formats structured prompts with context and instructions

03

LLM Processing

Ollama routes to quantized models with optimized inference

04

Response Generation

Transformer decodes tokens using positional embeddings and attention

05

Result Delivery

Structured JSON response parsed and displayed to user

Technology Stack

Backend & API

- Python 3.11 with type hints
- Flask for RESTful endpoints
- SQLAlchemy for data persistence
- WebSocket for streaming responses
- JWT for authentication

LLM Layer

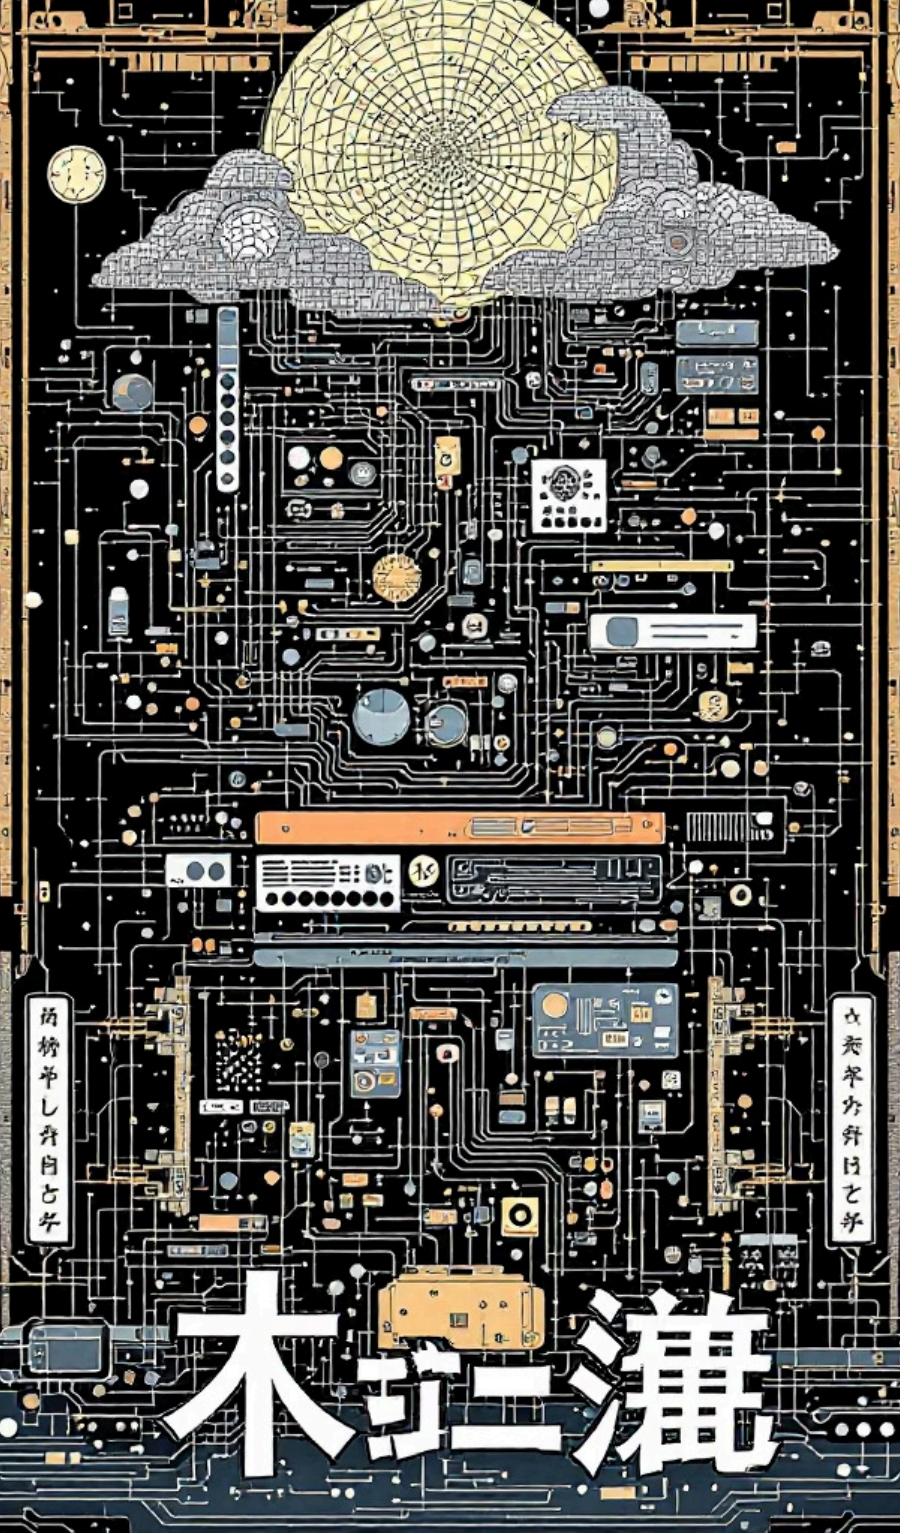
- Ollama for model management
- Llama3 8B/70B parameters
- Mistral 7B quantized variants
- GGUF format for efficiency

Frontend

- React with TypeScript
- CodeMirror for syntax highlighting
- WebSocket integration
- Material UI components

Infrastructure

- Docker containers
- NGINX reverse proxy
- Redis for caching
- PostgreSQL for metadata



LLM Architecture Deep Dive



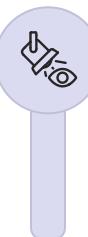
Tokenization

Byte-pair encoding splits code into subword units, mapping to vector embeddings



Transformer Layers

Multi-head self-attention computes token relationships using query, key, value projections



Context Processing

Positional encodings inject sequence order, enabling attention to weight relevant tokens



Token Prediction

Softmax output predicts next token probabilities, sampled during generation

Prompt Engineering Strategy

1

System Role Definition

Prepends "You are an expert software engineer" to establish persona and expertise level

2

Context Injection

Embeds surrounding code, error messages, and stack traces for situational awareness

3

Structured Output

Demands JSON format with specific fields: explanation, fix, complexity analysis

4

Chain-of-Thought

Requests step-by-step reasoning before final answer, improving accuracy and explainability

Backend formats prompts with context windows limited to model's token capacity (4K-32K tokens).

Model Selection & Optimization

Transformer Models

- Llama3 (8B, 7B): balanced performance
- Mistral 7B: faster inference
- Gemini variants: multimodal

Performance Metrics

- 2-3x faster inference vs. full-precision
- ~80% accuracy retention
- Real-time streaming

Quantization Benefits

- GGUF format reduces size 60%
- 4-bit quantization enables CPU-only
- Lower VRAM requirements

Latency Reduction

- Model caching in memory
- Prefill optimization
- Early stopping heuristics

Key Features



Code Explanation

Line-by-line breakdown of unfamiliar code with complexity analysis and pattern identification using AST parsing.



Optimization Engine

Identifies bottlenecks, suggests refactoring patterns, and provides Big-O analysis with alternative approaches.



Intelligent Debugging

Root cause analysis of errors with fix suggestions, stack trace correlation, and test case generation.



Interview Mode

DSA problem solving with step-by-step guidance, time/space complexity verification, and coding best practices.